

## ▼ Dog Breed Classification

In this project we will use traditional CNN, CNN with data augmentation and finally transfer Learning by VGG16 model with weights pre-on Imagenet to solve the dog breed classification problem

## ▼ Load Dataset Files

```
from google.colab import drive  
drive.mount('/content/drive')
```

☞ Go to this URL in a browser: [https://accounts.google.com/o/oauth2/auth?client\\_id=947318989803-6bn6qk8qdgf4n4g3pfee6491h](https://accounts.google.com/o/oauth2/auth?client_id=947318989803-6bn6qk8qdgf4n4g3pfee6491h)

```
Enter your authorization code:  
.....  
Mounted at /content/drive
```

Now, upload the given dataset file shared with you in your google drive and give its path for the below given `project_path` variable. For example, a path is given below according to the file path in our google drive. You need to change this to match the path of yours.

```
project_path = "/content/drive/My Drive/Conv-Project-2/"
```

```
import os  
  
os.chdir(project_path)
```

Run the below code to extract all the images in the train.zip files given in the dataset. We are going to use these images as train and va sets and their labels in further steps.

```
from zipfile import ZipFile  
with ZipFile(project_path+'train.zip', 'r') as z:  
    z.extractall()
```

Repeat the same step for test.zip

```
with ZipFile(project_path+'test.zip', 'r') as z:  
    z.extractall()
```

Repeat the same step for sample\_submission.csv.zip

```
from zipfile import ZipFile  
with ZipFile(project_path+'sample_submission.csv.zip', 'r') as z:  
    z.extractall()
```

Repeat the same step for labels.csv.zip

```
with ZipFile(project_path+'labels.csv.zip', 'r') as z:  
    z.extractall()
```

After this process, we will have 4 files - Train folder, test folder and labels.csv and sample\_submission.csv as part of your google drive

## ▼ Read labels.csv file using pandas

```
import pandas as pd
```

```
labels=pd.read_csv('labels.csv')
```

▼ Print the count of each category of Dogs given in the dataset

```
labels['breed'].value_counts()  
  
scottish_deerhound    126  
maltese_dog           117  
afghan_hound          116  
entlebucher            115  
bernese_mountain_dog  114  
...  
golden_retriever       67  
brabancon_griffon     67  
komondor               67  
briard                  66  
eskimo_dog              66  
Name: breed, Length: 120, dtype: int64
```

▼ Get one-hot encodings of labels

```
from sklearn.preprocessing import OneHotEncoder  
from sklearn.preprocessing import LabelEncoder  
  
label_encoder=LabelEncoder()  
  
labels_decoded=label_encoder.fit_transform(labels['breed'])  
  
%tensorflow_version 2.x  
import tensorflow as tf  
import numpy as np  
y = tf.keras.utils.to_categorical(labels_decoded, dtype='float32')  
y=np.array(y)
```

TensorFlow 2.x selected.

$y[0]$

## ▼ Preparing training dataset

1. Write a code which reads each and every id from labels.csv file and loads the corresponding image (in RGB - 128, 128, 3) from the train folder.
  2. Create 2 variables
    - a. x\_train - Should have all the images of the dogs from train folder
    - b. y\_train - Corresponding label of the dog

Note: The id of the dog images and its corresponding labels are available in labels.csv file

Hint: Watch the video shared on "Preparing the training dataset" if you face issue on creating the training dataset

```
img_rows=128  
img_cols=128  
num_channels=3
```

```
from tqdm import tqdm
import cv2
x_feature=[]
y_feature=[]


for dogid, img in tqdm(labels.values):
    train_img=cv2.imread('./train/{}.jpg'.format(dogid), 1)
    train_img=resize=cv2.resize(train_img, (img_rows, img_cols))
```

```
x_feature.append(train_img_resize)
y_feature.append(img)
```

```
↳ 100%|██████████| 10222/10222 [36:38<00:00, 5.52it/s]
```

```
import matplotlib.pyplot as plt
plt.imshow(x_feature[0])
```

```
↳ <matplotlib.image.AxesImage at 0x7fadf11689b0>
```



```
import numpy as np
```

```
x_feature=np.asarray(x_feature)
```

```
x_feature.shape
```

```
↳ (10222, 128, 128, 3)
```

Normalize the training data and convert into 4 dimensions so that it can be used as an input to conv layers in the model

```
X = x_feature.astype('float32')
```

[https://colab.research.google.com/drive/1kAZzo1gRF\\_NXHjatoEXnzoKL2aGxzejD#scrollTo=FM15LBzHJS1n&printMode=true](https://colab.research.google.com/drive/1kAZzo1gRF_NXHjatoEXnzoKL2aGxzejD#scrollTo=FM15LBzHJS1n&printMode=true)

```
X /= 255
```

```
X=X.reshape(X.shape[0],128,128,3)
```

```
X.shape
```

```
↳ (10222, 128, 128, 3)
```

## ▼ Split the training and validation data from x\_train\_data and y\_train\_data obtained from above step

```
from sklearn.model_selection import train_test_split
```

```
X_train_data , X_val, y_train_data, y_val = train_test_split(X, y, test_size=0.2, random_state=2)
print ("No. of images in train dataset: ", len(X_train_data))
print ("No. of images in Validation dataset: ", len(X_val))
```

```
↳ No. of images in train dataset: 8177
    No. of images in Validation dataset: 2045
```

```
X_train_data.shape
```

```
↳ (8177, 128, 128, 3)
```

## ▼ Loading the test data

Read the id column from the samples\_submission.csv and store it in test\_img

```
sample_sub=pd.read_csv('sample_submission.csv')
```

```
test_img=sample_sub['id']
```

```
import cv2
from tqdm import tqdm
```

Run the below code to load the test image files in x\_test\_feature


```
x_test_feature = []

for dogtestid in tqdm(test_img.values):
    test_img=cv2.imread('./test/{}.jpg'.format(dogtestid), 1)
    test_img_resize=cv2.resize(test_img, (img_rows, img_cols))
    x_test_feature.append(test_img_resize)
```

↳ 100% |██████████| 10357/10357 [37:12<00:00, 7.12it/s]

```
import matplotlib.pyplot as plt
plt.imshow(x_test_feature[0])
```

↳ <matplotlib.image.AxesImage at 0x7fca0c71f940>



Normalize the test data and convert it into 4 dimensions

```
x_test_feature=np.asarray(x_test_feature)
X_test= x_test_feature.astype('float32')
X_test /= 255
```

Build a basic conv neural network with 2 conv layers (kernel sizes - 5 and 3) add layers as mentioned below classification.

1. Add a Dense layer with 256 neurons with `relu` activation
2. Add a Dense layer with 120 neurons as final layer (as there are 120 classes in the given dataset) with `softmax` activation for classifiaction.

```
import tensorflow.keras
import keras
from tensorflow.keras.datasets import mnist
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Dropout, Flatten
from tensorflow.keras.layers import Conv2D, MaxPooling2D
input_shape = (img_rows, img_cols, 3)
```

↳ Using TensorFlow backend.

```
num_classes=120
#Initialize the model
model = Sequential()
#Add a Convolutional Layer with 32 filters of size 3X3 and activation function as 'ReLU'
model.add(Conv2D(32, kernel_size=(5, 5),
                 activation='relu',
                 input_shape=input_shape,name='conv_1'))
model.add(MaxPooling2D(pool_size=(2, 2),name='max_1', input_shape=(128,128,3)))

#Add a Convolutional Layer with 64 filters of size 3X3 and activation function as 'ReLU'
model.add(Conv2D(64, (3, 3), activation='relu',name='conv_2'))

#Add a MaxPooling Layer of size 2X2
model.add(MaxPooling2D(pool_size=(2, 2),name='max_2'))
```

```
model.add(Flatten())

model.add(Dense(256, activation='relu',name='dense_11'))

model.add(Dense(120, activation='softmax',name='dense_12'))

from tensorflow.keras.optimizers import SGD
from tensorflow.keras.losses import categorical_crossentropy

#To use adam optimizer for learning weights with learning rate = 0.001
optimizer = SGD(lr=0.001)
#Set the loss function and optimizer for the model training
model.compile(loss=categorical_crossentropy,optimizer=optimizer,metrics=['accuracy'])

model.summary()
```

↳ Model: "sequential\_1"

Layer (type)	Output Shape	Param #
<hr/>		
conv_1 (Conv2D)	(None, 124, 124, 32)	2432
max_1 (MaxPooling2D)	(None, 62, 62, 32)	0
conv_2 (Conv2D)	(None, 60, 60, 64)	18496
max_2 (MaxPooling2D)	(None, 30, 30, 64)	0
flatten_1 (Flatten)	(None, 57600)	0
dense_11 (Dense)	(None, 256)	14745856
dense_12 (Dense)	(None, 120)	30840
<hr/>		
Total params: 14,797,624		
Trainable params: 14,797,624		
Non-trainable params: 0		

▼ Use batch\_size = 128 and epochs = 10 and execute the model

```
batch_size=128
epoches=10
model.fit(X_train_data, y_train_data,
           batch_size=batch_size,
           epoches=epoches,
           verbose=1,
           validation_data=(X_val, y_val))
```

↳ Train on 8177 samples, validate on 2045 samples

```
Epoch 1/10
8177/8177 [=====] - 4s 548us/sample - loss: 4.7566 - accuracy: 0.0212 - val_loss: 4.7746 - val
Epoch 2/10
8177/8177 [=====] - 4s 539us/sample - loss: 4.7535 - accuracy: 0.0204 - val_loss: 4.7728 - val
Epoch 3/10
8177/8177 [=====] - 4s 534us/sample - loss: 4.7497 - accuracy: 0.0218 - val_loss: 4.7702 - val
Epoch 4/10
8177/8177 [=====] - 4s 534us/sample - loss: 4.7457 - accuracy: 0.0207 - val_loss: 4.7676 - val
Epoch 5/10
8177/8177 [=====] - 4s 539us/sample - loss: 4.7413 - accuracy: 0.0219 - val_loss: 4.7645 - val
Epoch 6/10
8177/8177 [=====] - 4s 540us/sample - loss: 4.7364 - accuracy: 0.0234 - val_loss: 4.7617 - val
Epoch 7/10
8177/8177 [=====] - 4s 543us/sample - loss: 4.7309 - accuracy: 0.0225 - val_loss: 4.7581 - val
Epoch 8/10
8177/8177 [=====] - 4s 540us/sample - loss: 4.7254 - accuracy: 0.0227 - val_loss: 4.7546 - val
Epoch 9/10
8177/8177 [=====] - 4s 537us/sample - loss: 4.7190 - accuracy: 0.0240 - val_loss: 4.7517 - val
Epoch 10/10
8177/8177 [=====] - 4s 541us/sample - loss: 4.7121 - accuracy: 0.0246 - val_loss: 4.7458 - val
<tensorflow.python.keras.callbacks.History at 0x7fadca3f2390>
```

```
score = model.evaluate(X_val, y_val)
print('Test loss:', score[0])
print('Test accuracy:', score[1])
```

↳

```
2045/2045 [=====] - 1s 287us/sample - loss: 4.7458 - accuracy: 0.0181
Test loss: 4.745807010911788
Test accuracy: 0.01809291
```

- ▼ The model accuracy is very poor !!!!
- ▼ Use Data Augmentation in the above model to see if the accuracy improves

```
from tensorflow.keras.preprocessing.image import img_to_array,load_img,ImageDataGenerator
from tensorflow.keras.layers import ZeroPadding2D

train_datagen=ImageDataGenerator(
    rotation_range=30,
    width_shift_range=0.2,
    height_shift_range=0.2,
    rescale=1./255,
    shear_range=0.2,
    zoom_range=0.2,
    horizontal_flip=True,
    fill_mode='nearest'
)
```

- ▼ Using the above objects, create the image generators with variable names `train_generator` and `val_generator`.  
You need to use `train_datagen.flow()` and `val_datagen.flow()`

```
train_generator=train_datagen.flow(X_train_data, y=y_train_data, batch_size=32)
val_generator=train_datagen.flow(X_val, y=y_val, batch_size=32)
```

- ▼ Fit the model using `fit_generator()` using `train_generator` and `val_generator` from the above step with 10 epochs.

```
model.fit_generator(  
    train_generator,  
    epochs=10,  
    validation_data=val_generator)
```



## ▼ Model accuracy is still poor!!!

## ▼ Lets use Transfer Learning

Download the vgg wieght file from here :

[https://github.com/MinerKasch/applied\\_deep\\_learning/blob/master/vgg16\\_weights\\_tf\\_dim\\_ordering\\_tf\\_kernels\\_notop.h5](https://github.com/MinerKasch/applied_deep_learning/blob/master/vgg16_weights_tf_dim_ordering_tf_kernels_notop.h5)

Train for 256 steps validate for 64 steps

Use the below code to load VGG16 weights trained on ImageNet

Epoch 2/10

```
from tensorflow.keras.applications.vgg16 import VGG16, preprocess_input
# Instantiate the model with the pre-trained weights (no top)
base_model= VGG16(weights=(project_path+'vgg16_weights_tf_dim_ordering_tf_kernels_notop.h5'),
                  include_top=False, pooling='avg')
```

Epoch 5/10

256/256 [

100% 140ms/step - loss: 4.7054 - accuracy: 0.0125 - val\_loss: 4.7054 - val\_accuracy: 0.0125

Print the summary of the base\_model

Epoch 8/10

```
base_model.summary()
```



Model: "vgg16"

Layer (type)	Output Shape	Param #
<hr/>		
input_2 (InputLayer)	[(None, None, None, 3)]	0
block1_conv1 (Conv2D)	(None, None, None, 64)	1792
block1_conv2 (Conv2D)	(None, None, None, 64)	36928
block1_pool (MaxPooling2D)	(None, None, None, 64)	0
block2_conv1 (Conv2D)	(None, None, None, 128)	73856
block2_conv2 (Conv2D)	(None, None, None, 128)	147584
block2_pool (MaxPooling2D)	(None, None, None, 128)	0
block3_conv1 (Conv2D)	(None, None, None, 256)	295168
block3_conv2 (Conv2D)	(None, None, None, 256)	590080
block3_conv3 (Conv2D)	(None, None, None, 256)	590080
block3_pool (MaxPooling2D)	(None, None, None, 256)	0
block4_conv1 (Conv2D)	(None, None, None, 512)	1180160
block4_conv2 (Conv2D)	(None, None, None, 512)	2359808
block4_conv3 (Conv2D)	(None, None, None, 512)	2359808
block4_pool (MaxPooling2D)	(None, None, None, 512)	0
block5_conv1 (Conv2D)	(None, None, None, 512)	2359808
block5_conv2 (Conv2D)	(None, None, None, 512)	2359808
block5_conv3 (Conv2D)	(None, None, None, 512)	2359808
block5_pool (MaxPooling2D)	(None, None, None, 512)	0
global_average_pooling2d_1 (	(None, 512)	0

```
=====
Total params: 14,714,688
Trainable params: 14,714,688
Non-trainable params: 0
```

---

▼ Add the following classification layers to the imported VGG Model

1. Flatten Layer
2. Dense layer with 1024 neurons with activation as Relu
3. Dense layer with 256 neurons with activation as Relu
4. Dense layer with 120 neurons with activation as Softmax

```
from tensorflow.keras.models import Model

x=base_model.output
x=Flatten()(x)
x=Dense(1024,activation='relu')(x) #we add dense layers so that the model can learn more complex functions and classify for
x=Dense(256,activation='relu')(x) #dense layer 2
preds=Dense(120,activation='softmax')(x) #final layer with softmax activation

vgmodel=Model(inputs=base_model.input,outputs=preds)

for i,layer in enumerate(vgmodel.layers):
    print(i,layer.name)
```



```
0 input_2  
1 block1_conv1  
2 block1_conv2  
3 block1_pool  
4 block2_conv1  
5 block2_conv2  
6 block2_pool  
7 block3_conv1  
8 block3_conv2  
9 block3_conv3  
10 block3_pool  
11 block4_conv1  
12 block4_conv2  
13 block4_conv3  
14 block4_pool  
15 block5_conv1  
16 block5_conv2  
17 block5_conv3  
18 block5_pool  
19 global_average_pooling2d_1  
20 flatten_4  
21 dense_6  
22 dense_7  
23 dense_8
```

```
#To use adam optimizer for learning weights with learning rate = 0.001  
optimizer = SGD(lr=0.001)  
#Set the loss function and optimizer for the model training  
vgmodel.compile(loss=categorical_crossentropy,optimizer=optimizer,metrics=[ 'accuracy' ])
```

## ▼ Make all the layers in the base\_model (VGG16) to be non-trainable

```
for layer in vgmodel.layers:  
    if('dense' not in layer.name): #prefix detection to freeze layers which does not have dense  
        #Freezing a layer  
        layer.trainable = False
```

▼ Fit and compile the model with batch\_size = 128 and epochs = 10 and execute the model

Try to get training and validation accuracy to be more than 90%

```
batch_size=128  
epoches=10  
vgmodel.fit(train_generator,  
            epochs=epoches,  
            verbose=1,  
            validation_data=(val_generator))
```



```
WARNING:tensorflow:sample_weight modes were coerced from
...
    to
    [...]
WARNING:tensorflow:sample_weight modes were coerced from
...
    to
    [...]
Train for 256 steps, validate for 64 steps
Epoch 1/10
256/256 [=====] - 57s 225ms/step - loss: 4.7908 - accuracy: 0.0108 - val_loss: 4.7945 - val_ac
Epoch 2/10
256/256 [=====] - 49s 193ms/step - loss: 4.7854 - accuracy: 0.0110 - val_loss: 4.7935 - val_ac
Epoch 3/10
256/256 [=====] - 50s 194ms/step - loss: 4.7829 - accuracy: 0.0105 - val_loss: 4.7925 - val_ac
Epoch 4/10
256/256 [=====] - 49s 193ms/step - loss: 4.7816 - accuracy: 0.0108 - val_loss: 4.7914 - val_ac
Epoch 5/10
256/256 [=====] - 50s 194ms/step - loss: 4.7808 - accuracy: 0.0108 - val_loss: 4.7919 - val_ac
Epoch 6/10
256/256 [=====] - 49s 192ms/step - loss: 4.7800 - accuracy: 0.0115 - val_loss: 4.7915 - val_ac
Epoch 7/10
256/256 [=====] - 50s 194ms/step - loss: 4.7797 - accuracy: 0.0117 - val_loss: 4.7917 - val_ac
Epoch 8/10
256/256 [=====] - 50s 195ms/step - loss: 4.7793 - accuracy: 0.0122 - val_loss: 4.7919 - val_ac
Epoch 9/10
256/256 [=====] - 50s 194ms/step - loss: 4.7789 - accuracy: 0.0109 - val_loss: 4.7920 - val_ac
Epoch 10/10
256/256 [=====] - 50s 194ms/step - loss: 4.7785 - accuracy: 0.0113 - val_loss: 4.7916 - val_ac
<tensorflow.python.keras.callbacks.History at 0x7fadcc25f9b00>
```

